

# Scheduling

# Main Points

រៀបចំវិធានថា thread ដីមកទៀត៖ run នៅ

- Scheduling policy: what to do next, when there are multiple threads ready to run
  - Or multiple packets to send, or web requests to serve, or ...
- Definitions ធម្ម័យ (ព័ត៌មានមិនមែន) – response time, throughput, predictability
- Uniprocessor policies វិធានការងារមួយគោល ឬ អេដូលុយ  
– FIFO, round robin, optimal
  - multilevel feedback as approximation of optimal
- Multiprocessor policies វិធានមួយពីរ ឬ ពីច្បាស់
  - Affinity scheduling, gang scheduling
- Queueing theory ការងារក្នុងការគិត
  - Can you predict/improve a system's response time?

# Definitions

- Task/Job
  - User request: e.g., mouse click, web request, shell command, ...
- Latency/response time
  - How long does a task take to complete? ឧបនីមួយៗនេះ
- Throughput ការសែវភ័យ [Throughput ឬបានដល់រៀបចំនៃការប្រឡង]
  - How many tasks can be done per unit of time?
- Overhead
  - How much extra work is done by the scheduler?
- Fairness គ្មានការកំណត់នៅមិនមែន  
  - How equal is the performance received by different users?
- Predictability ការកំណត់នៅមិនមែន  
  - How consistent is the performance over time?

ការកំណត់នៅមិនមែន មិនមែនការកំណត់នៃការប្រឡង

មិន

ជាប្រព័ន្ធដែលមិនមែន  
ការកំណត់នៅមិនមែន

# More Definitions

- Workload **ក្រោមងារ នៅលើបានយក នៃការធ្វើដំឡើ**
  - Set of tasks for system to perform
- Preemptive scheduler **ដែលត្រូវការយកចុចការឡើង**
  - If we can take resources away from a running task
- Work-conserving **ការរួមទៅនឹង នៃការ ដែលមានលក្ខណៈថា មិនអាចបានយកចុចការឡើង**
  - Resource is used whenever there is a task to run
- Scheduling algorithm **ដែលបានបង្កើតឡើង**
  - takes a workload as input ③ ដែលបានបង្កើតឡើង
  - decides which tasks to do first គោរពដែល ឱ្យដែលការិកនូវ workload ដែលមានចំណាំ ការងាររបស់ 4 processor
  - Performance metric (throughput, latency) as output ផ្តល់ជាផ្លូវការ
  - Only preemptive, work-conserving schedulers to be considered ត្រូវបានគេចាប់អាមេរី

# First In First Out (FIFO)

- Schedule tasks in the order they arrive
  - Continue running them until they complete or give up the processor
- Example: memcached
  - Facebook cache of friend lists, ...
- On what workloads is FIFO particularly bad?

ការងារ = តីវេរក់របស់ខ្លួន

ទំនួន workload

ពេលវេលាការណ៍

# Shortest Job First (SJF)

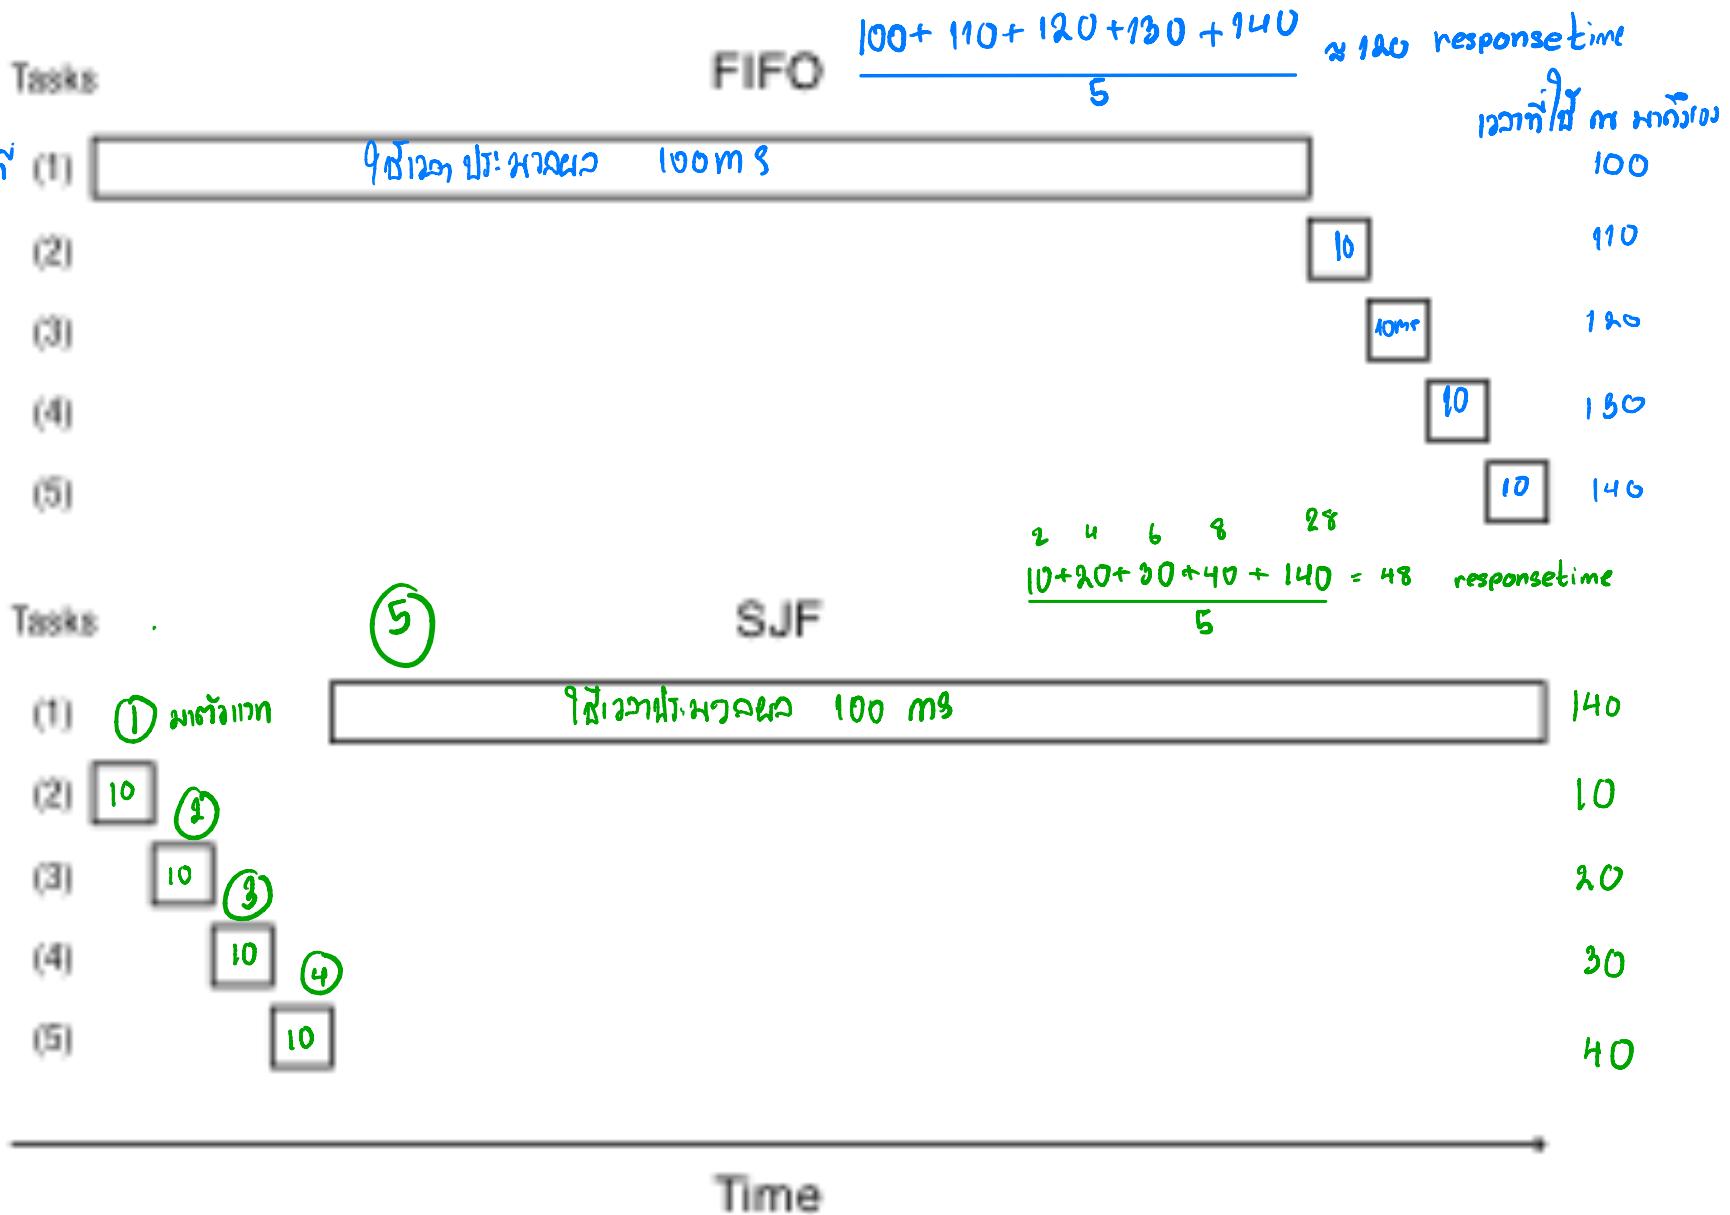
ការងារដែលត្រូវបានធ្វើ = ការងារដែលត្រូវបានធ្វើ

- ការងារដែលត្រូវបានធ្វើមេត្តុរវាង គ្មាន Job បានអនុញ្ញាត - ឬ

- កិច្ចការ || ចងក់

- Always do the task that has the shortest remaining amount of work to do
  - Often called Shortest Remaining Time First (SRTF)
- Suppose we have five tasks arrive one right after each other, but the first one is much longer than the others
  - Which completes first in FIFO? Next?
  - Which completes first in SJF? Next?

# FIFO vs. SJF



# Question

- Claim: SJF is optimal for average response time

– Why? กាំង response time អត្ថប័ណ្ណអិនមេ

ទេវារៀនអិនមេ ការគោរព: ពាំ

- Does SJF have any downsides? អវត្ថុភ័យ
- អាកាសក្រឡាកំវិនាក់  
( មែន: ដែល ធ្វើនា ឬ workload ) ដើម្បី សារ៖ "starvation" នាយករា

res time = response time

# Question

ສິ້ນສາພາກ ອີນຕ່າງລະບົບທີ່ຕັ້ງຮູ້ໃນ

- Is FIFO ever optimal? ສໍາຜົນ ກາຕື່ລົບ

ມີ ຖໍາ Job ເທົ່ານາແບບເຊຍວໍາດັບ ປົດຍາວິທະຍາກົມ-ສາກ

FIFO res time = SJF res time

ສິ້ນສາພາກ ອີນຕ່າງລະບົບທີ່ແຈກຕື່ອນຫຼືຂຶ້ນ

- Pessimal? ສໍາຜົນ ກາຕື່ລົບ

ມີ ກາຕື່ Job ມາຮັດກີ່ສຸດ - ນັບຍັງສຸດ

កិច្ចការទែរការ / ពាណិជ្ជការ  
ជាដៃលើការងារ

# Round Robin

- Each task gets resource for a fixed period of time (time quantum) (មាត្រាកំណើន និងរយៈពេល)
  - If task doesn't complete, it goes back in line
- Need to pick a time quantum
  - What if time quantum is too long? ស្ថិតិ?
    - Infinite? Job ការងារនេះ តាមការណែនាំ និងការបញ្ចូលការងារ ដូចតួនាទីនេះ
  - What if time quantum is too short? ធ្លើ?
    - One instruction? តាមការណែនាំ និងការបញ្ចូលការងារ ដូចតួនាទីនេះ Overhead ខ្សោយការ

ສົ່ງມາເປັນດີປະເງິນ

SJF

ຕົວທາງການ ການກຳຈະ Job

ໃນກຳນົດ time quantum 9 ນັ້ນກະລຸງ

# Round Robin

time quantum 1ms

Tasks

ຕົ່ງດົງກົວດັບ

Round Robin (1 ms time slice)

(1)

(2)

(3)

(4)

(5)

Rest of Task 1

...

Tasks

Round Robin (100 ms time slice)

104 + ການກຳ process

(1)

100

Rest of Task 1

(2)

101

(3)

102

(4)

103

(5)

104

Time

# Round Robin vs. FIFO

మాన్‌గారు

టాండమ్‌<sup>స్ట</sup>

- Assuming zero-cost time slice, is Round Robin always better than FIFO?

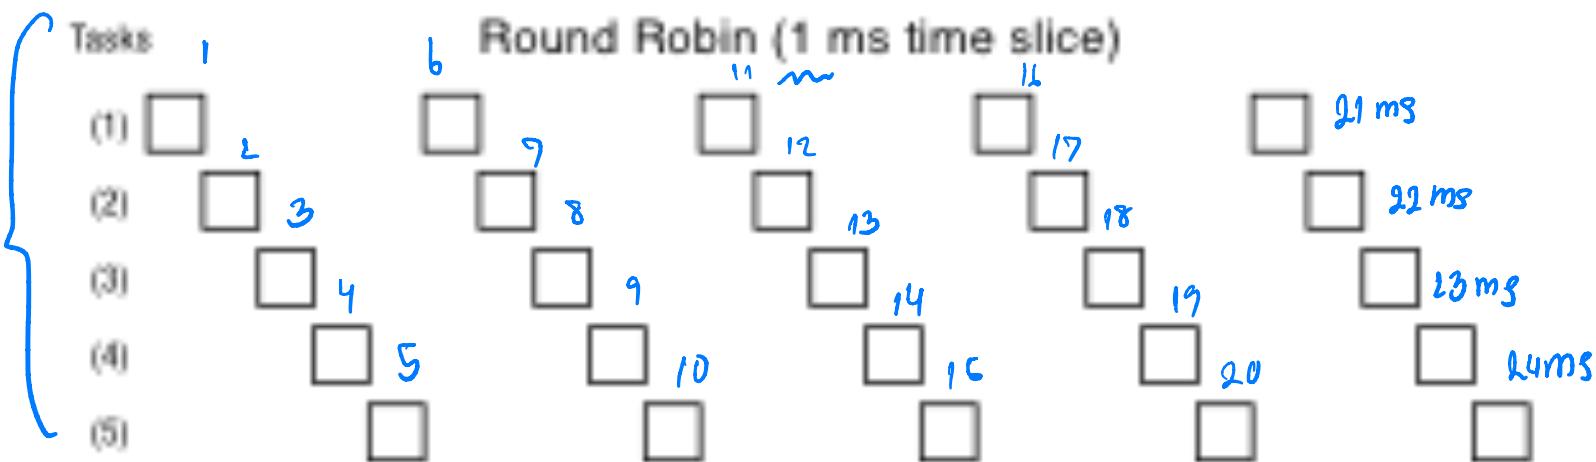
మాన్‌గారు →

1 Job = 5 ms

# Round Robin vs. FIFO

Job in Slice 0 until 5 slice

กิจกรรม  
\_FIFO



Tasks

FIFO and SJF



Time

# Round Robin = Fairness?

SJF เช่นเดียว

- Is Round Robin always fair?

ถ้า FIFO ก็เท่ากัน 1 time quantum ความก่อตัวใช้เวลา

- What is fair?

- FIFO? ความก่อตัวใช้เวลา resource ก่อน

- Equal share of the CPU? แนวหน้าตัดลง CPU ทุกๆ เท่านั้น

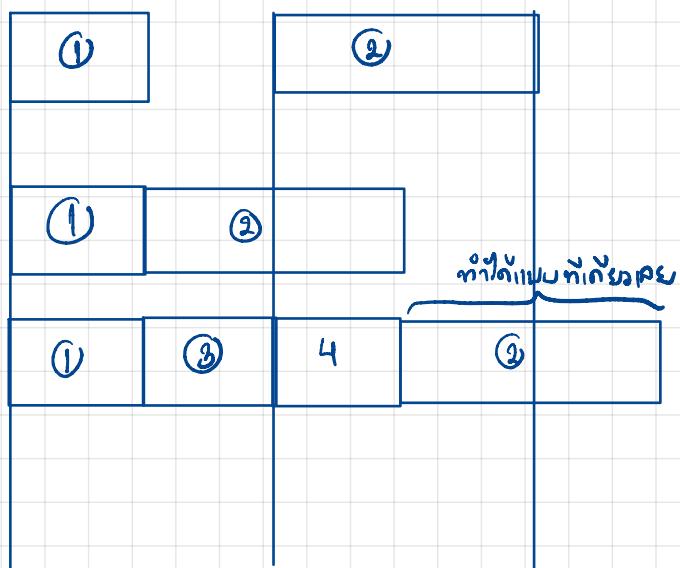
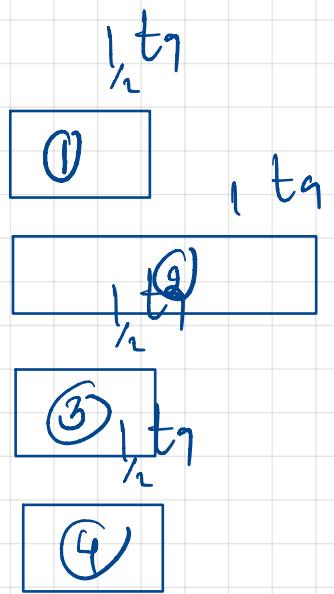
- What if some tasks don't need their full share?  
ก่อตัวไม่เต็ม 1 time quantum ผู้ใช้ทรัพยากรน้อย Round Robin ไม่ดี

- Minimize worst case divergence?  
ก่อตัวต้องสั้นกว่า 1 Job ต่อ time quantum ขึ้นไป (สูง กี่ ขนาด 1 ที่ ก่อตัวในพื้นที่มากกว่าตัวเดียว)

- Time task would take if no one else was running

- Time task takes under scheduling algorithm  
เวลาต่อๆ กันมา ไม่ต่อเนื่อง หมายความ = ว่า ไม่แน่นอน = ไม่ต่อเนื่อง

- ขนาด กี่ ก่อตัว ไม่ต่อเนื่อง หมายความ = ว่า ไม่แน่นอน = ไม่ต่อเนื่อง ? ก่อตัวไม่ต่อเนื่อง

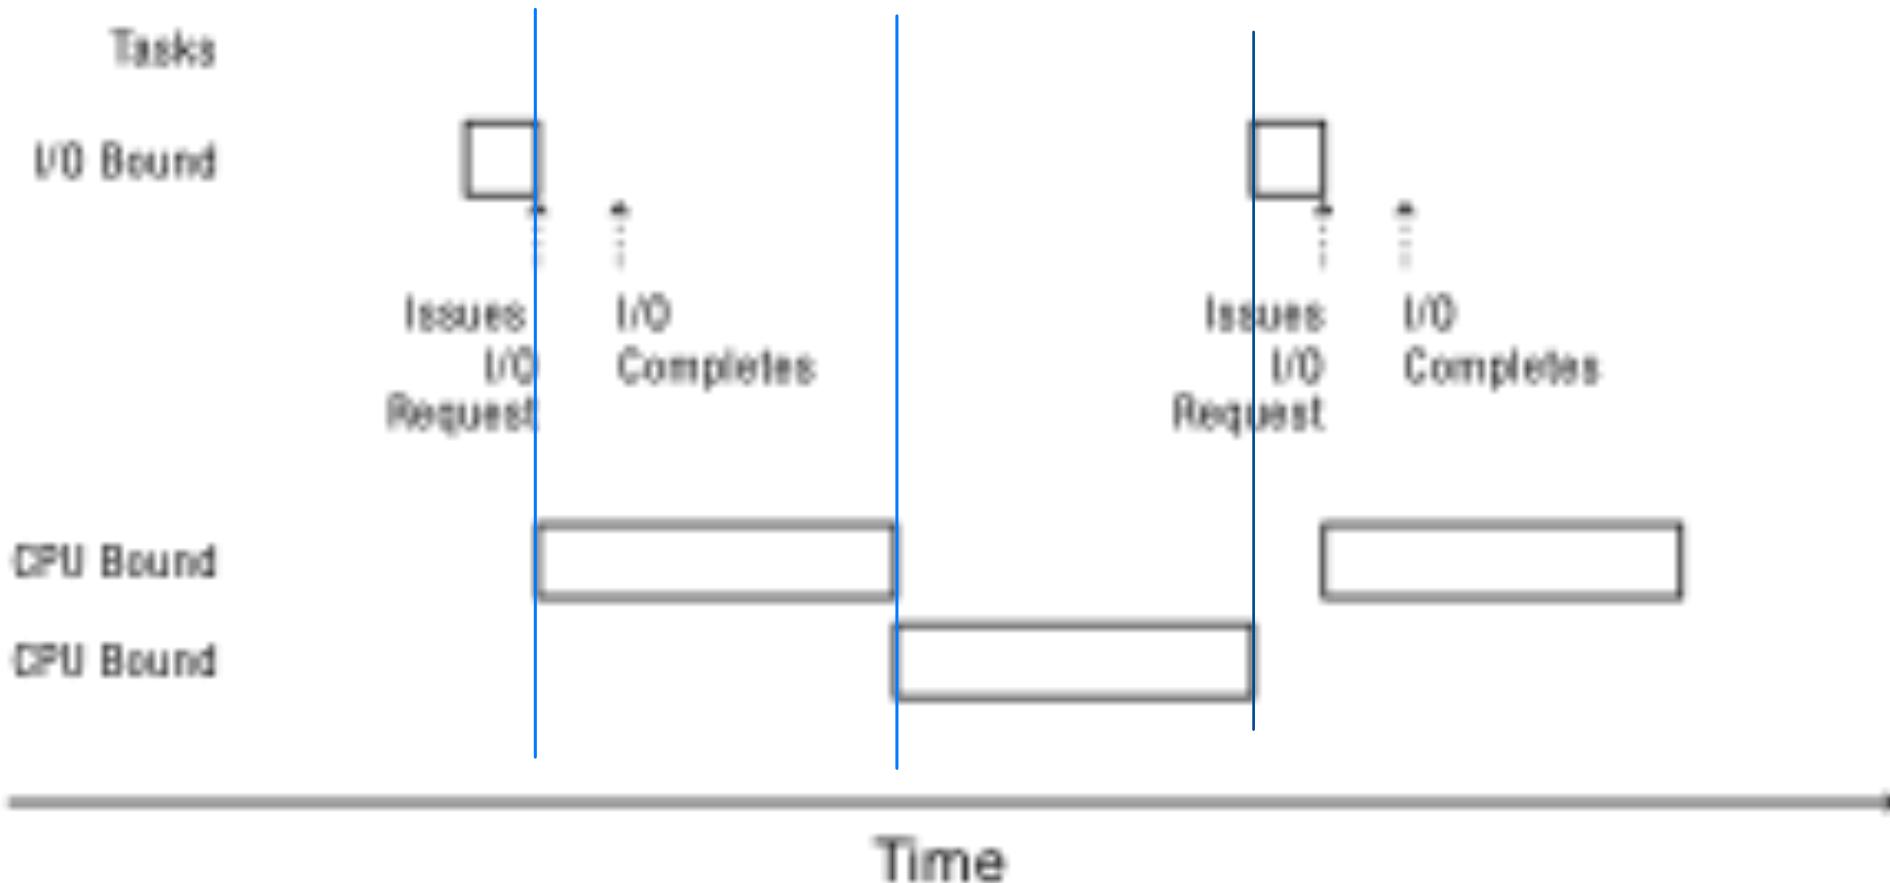


ມີເຮັດວຽກນັ້ນ  
ມີຮັດວຽກນັ້ນ  
ມີຮັດວຽກນັ້ນ

RR Fairness

ຕົກລາງລູກໆນັ້ນຕົກລາງລູກໆນັ້ນ

# Mixed Workload



ເລືດ ຈົບ ກໍ່ມີເວວາ ໄກສູ່ໃນຈະນຸ່ງ

# Max-Min Fairness

- ອິນເຕັມ ເຮັດວຽກ  
- ຖັນຍາ ຂອບທິປະໄຕ ໃຫຍວດການຄະນະ ແລ້ວ

WOW!

- How do we balance a mixture of repeating tasks:
  - Some I/O bound, need only a little CPU
  - Some compute bound, can use as much CPU as they are assigned
- One approach: maximize the minimum allocation given to a task
  - If any task needs less than an equal share, schedule the smallest of these first
  - Split the remaining time using max-min
  - If all remaining tasks need at least equal share, split evenly

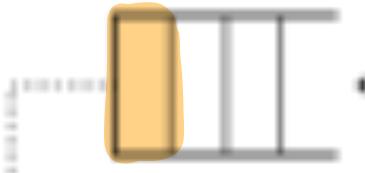
# Multi-level Feedback Queue (MFQ)

- Goals:
  - Responsiveness ตอบสนับเร็ว
  - Low overhead
  - Starvation freedom ไม่ถูก忽落
  - Some tasks are high/low priority
    - หา  $\max\text{-}\min$
  - Fairness (among equal priority tasks)
- Not perfect at any of them!
  - Used in Linux (ระบบ and probably Windows, MacOS) ใช้ใน Linux

# MFQ

- Set of Round Robin queues
  - Each queue has a separate priority
- High priority queues have short time slices
  - Low priority queues have long time slices
- Scheduler picks first thread in highest priority queue
- Tasks start in highest priority queue
  - If time slice expires, task drops one level

# MFQ

Priority	Time Slice (ms)	Round Robin Queues	
1	10		New or I/O Bound Task
2	20		Time Slice Expiration
3	40		
4	80		

- ถ้า Job 9 ใน Queue 1 หมด : ไปที่ Job 9 ใน Queue 2  
 - ทำงานไม่ตอบโต้แล้ว 2 8 ต่อ priority  
 เนื่องจาก / tq จะได้มากขึ้น

- ก้าทำงานใน tq ก็ต้อง Queue เดิม + ก้าตรวจสอบการทำงาน  
 - เมื่อ 9 ต่อ ถ้าไม่ทำงานต่อ (timeout) ; จะส่งไป Queue 2  
 \* ไม่ให้หักก็ต้องเรียบ ก่อนย้าย หาอยู่ Queue 1

# Uniprocessor Summary (1)

- FIFO is simple and minimizes overhead.
- If tasks are variable in size, then FIFO can have very poor average response time.
- If tasks are equal in size, FIFO is optimal in terms of average response time.
- Considering only the processor, SJF is optimal in terms of average response time.
- SJF is <sup>optimal</sup> <sub>pessimal</sub> in terms of variance in response time.

# Uniprocessor Summary (2)

Job պատճենագործ

- If tasks are variable in size, Round Robin approximates SJF.

Job պատճենագործ - պահանջ

- If tasks are equal in size, Round Robin will have very poor average response time.

Job գովարժականացնելու ընթացքում max-min ռեզուլտատ

- Tasks that intermix processor and I/O benefit from SJF and can do poorly under Round Robin.

# Uniprocessor Summary (3)

- Max-Min fairness can improve response time for I/O-bound tasks.
- Round Robin and Max-Min fairness both avoid starvation.  
ພວກເຮົາ ຕ່າງ , ຕົວທີ່ຢູ່ , time out , priority ສໍາກຳນາດຂອງ
- By manipulating the assignment of tasks to priority queues, an MFQ scheduler can achieve a balance between responsiveness, low overhead, and fairness.  
ແນວໃຈກັບ  
RR + SJF